ALGORITHM & DOCUMENTATION: MINRES-QLP for Singular Symmetric and Hermitian Linear Equations and Least-Squares Problems

SOU-CHENG T. CHOI
University of Chicago/Argonne National Laboratory
and
MICHAEL A. SAUNDERS
Stanford University

We describe algorithm MINRES-QLP and its FORTRAN 90 implementation for solving symmetric or Hermitian linear systems or least-squares problems. If the system is singular, MINRES-QLP computes the unique minimum-length solution (also known as the pseudoinverse solution), which generally eludes MINRES. In all cases, it overcomes a potential instability in the original MINRES algorithm. A positive-definite preconditioner may be supplied. Our FORTRAN 90 implementation illustrates a design pattern that allows users to make problem data known to the solver but hidden and secure from other program units. In particular, we circumvent the need for reverse communication. Example test programs input and solve real or complex problems specified in Matrix Market format. While we focus here on a FORTRAN 90 implementation, we also provide and maintain MATLAB versions of MINRES and MINRES-QLP.

---

1. INTRODUCTION

MINRES-QLP [Choi 2006; Choi et al. 2011] is a Krylov subspace method for computing the minimum-length and minimum-residual solution (also known as the pseudoinverse solution) x to the following linear systems or least-squares (LS) problems:

$$\text{solve } Ax = b, \tag{1}$$

$$\text{minimize } \|x\|_2 \quad \text{s.t.} \quad Ax = b, \tag{2}$$

$$\text{minimize } \|x\|_2 \quad \text{s.t.} \quad x \in \arg\min_x \|Ax - b\|_2, \tag{3}$$

where A is an $n \times n$ symmetric or Hermitian matrix and b is a real or complex n-vector. Problems (1) and (2) are treated as special cases of (3). The matrix A is usually large and sparse, and it may be singular.[1] It is defined by means of a user-written subroutine `Aprod`, whose function is to compute the product $y = Av$ for any given vector v.

Let x_k be the solution estimate associated with MINRES-QLP's kth iteration, with residual vector $r_k = b - Ax_k$. Without loss of generality, we define $x_0 = 0$. MINRES-QLP provides recurrent estimates of $\|x_k\|$, $\|r_k\|$, $\|Ar_k\|$, $\|A\|$, $\text{cond}(A)$, and $\|Ax_k\|$, which are used in the stopping conditions.

Other iterative methods specialized for symmetric systems $Ax = b$ are the conjugate-gradient method (CG) [Hestenes and Stiefel 1952], SYMMLQ and MINRES [Paige and Saunders 1975], and SQMR [Freund and Nachtigal 1994]. Each method requires one product Av_k at each iteration for some vector v_k. CG is intended for positive-definite A, whereas the other solvers allow A to be indefinite.

If A is singular, SYMMLQ requires the system to be consistent, whereas MINRES returns an LS solution for (3) but generally not the min-length solution; see [Choi 2006; Choi et al. 2011] for examples. SQMR without preconditioning is mathematically equivalent to MINRES but could fail on a singular problem. To date, MINRES-QLP is probably the most suitable CG-type method for solving (3).

In some cases the more established symmetric methods may still be preferable.

(1) If A is positive definite, CG minimizes the energy norm of the error $\|x - x_k\|_A$ in each Krylov subspace and requires slightly less work per iteration. However, CG, MINRES, and MINRES-QLP do reduce $\|x - x_k\|_A$ and $\|x - x_k\|$ monotonically. Also, MINRES and MINRES-QLP often reduce $\|r_k\|$ to the desired level significantly sooner than does CG, and the backward error for each x_k decreases monotonically. (See Section 2.4 and [Fong 2011; Fong and Saunders 2012].)

(2) If A is indefinite but $Ax = b$ is consistent (e.g., if A is nonsingular), SYMMLQ requires slightly less work per iteration, and it reduces the error norm $\|x - x_k\|$ monotonically. MINRES and MINRES-QLP *usually* reduce $\|x - x_k\|$ [Fong 2011; Fong and Saunders 2012].

(3) If A is indefinite and well-conditioned and $Ax = b$ is consistent, MINRES might be preferable to MINRES-QLP because it requires the same number of iterations but slightly less work per iteration.

---

[1] A further input parameter σ (a real shift parameter) causes MINRES-QLP to treat "A" as if it were $A - \sigma I$. For example, "singular A" really means that $A - \sigma I$ is singular.

(4) MINRES and MINRES-QLP require a preconditioner to be positive definite. SQMR might be preferred if A is indefinite and an effective indefinite preconditioner is available.

MINRES-QLP has two phases. Iterations start in the *MINRES phase* and transfer to the *MINRES-QLP phase* when a subproblem (see (8) below) becomes ill-conditioned by a certain measure. If every subproblem is of full rank and well-conditioned, the problem can be solved entirely in the MINRES phase, where the cost per iteration is essentially the same as for MINRES. In the MINRES-QLP phase, one more work vector and $5n$ more multiplications are used per iteration.

MINRES-QLP described here is implemented in FORTRAN 90 for real double-precision problems. It contains no machine-dependent constants and does not need to use features such as polymorphism from FORTRAN 2003 or 2008. It requires an auxiliary subroutine `Aprod` and, if a preconditioner is supplied, a second subroutine `Msolve`. We also provide a complex implementation for Hermitian problems. Precision other than double can be obtained by changing one line of code. The programs can be compiled with FORTRAN 90 and FORTRAN 95 compilers such as `g95` and `gfortran`.

We also maintain a MATLAB implementation capable of solving both real and complex problems. All implementations are available at [SOL] or the first author's homepage, http://home.uchicago.edu/sctchoi/.

Table I lists the main notation used.

Table I.   Key notation.

$\| \cdot \|$	matrix or vector two-norm
$\bar{A}$	$\bar{A} = A - \sigma I$ (see also σ below)
$\mathrm{cond}(A)$	condition number of A with respect to two-norm $= \frac{\max \vert \lambda_i \vert}{\min_{\lambda_i \neq 0} \vert \lambda_i \vert}$
e_i	ith unit vector
ℓ	index of the last Lanczos iteration when $\beta_{\ell+1} = 0$
n	order of A
$\mathrm{null}(A)$	null space of A defined as $\{x \in \mathbb{R}^n \mid Ax = 0\}$
$\mathrm{range}(A)$	column space of A defined as $\{Ax \mid x \in \mathbb{R}^n\}$
T	(right superscript to a vector or a matrix) transpose
$x^\dagger$	unique minimum-length least-squares solution of problem (3)
$\mathcal{K}_k(A, b)$	kth Krylov subspace defined as span$\{b, Ab, \ldots, A^{k-1}b\}$
ε	machine precision
σ	real scalar shift to diagonal of A

1.1   Least-Squares Methods

Further existing methods that could be applied to (3) are CGLS and LSQR [Paige and Saunders 1982a; Paige and Saunders 1982b], LSMR [Fong and Saunders 2011], and GMRES [Saad and Schultz 1986], all of which reduce $\|r_k\|$ monotonically. The first three methods would require two products Av_k and Au_k each iteration and would be generating points in less favorable subspaces. GMRES requires only products Av_k and could use any nonsingular (possibly indefinite) preconditioner. It needs increasing storage and work each iteration, perhaps requiring restarts, but it

Table II. Comparison of various least-squares solvers on $n \times n$ systems (3). Storage refers to memory required by working vectors in the solvers. Work counts number of floating-point multiplications. On inconsistent systems, all solvers below except MINRES and GMRES with restart parameter m return the minimum-length LS solution (assuming no preconditioner).

Solver	Storage	Work per Iteration	Products per Iteration	Systems to Solve per Iteration with Preconditioner
MINRES	$7n$	$9n$	1	1
MINRES-QLP	$7n$–$8n$	$9n$–$14n$	1	1
GMRES(m)	$(m+2)n$	$(m+3+1/m)n$	1	1
CGLS	$4n$	$5n$	2	2
LSQR	$5n$	$8n$	2	2
LSMR	$6n$	$9n$	2	2

could be more effective than MINRES or MINRES-QLP (and the other solvers) if few total iterations were required. Table II summarizes the computational requirements of each method.

1.2   Regularization

We do not discourage using CGLS, LSQR, or LSMR if the goal is to regularize an ill-posed problem using a small damping factor $\lambda > 0$ as follows:

$$\min_x \left\| \begin{bmatrix} A \\ \lambda I \end{bmatrix} x - \begin{bmatrix} b \\ 0 \end{bmatrix} \right\|. \tag{4}$$

However, this approach destroys the original problem's symmetry. The normal equation of (4) is $(A^2 + \lambda^2 I)x = Ab$, which suggests that a diagonal shift to A may well serve the same purpose in some cases. For symmetric positive-definite A, $\bar{A} = A - \sigma I$ with $\sigma < 0$ enjoys a smaller condition number. When A is indefinite, a good choice of σ may not exist, for example, if the eigenvalues of A were symmetrically positioned around zero. When this symmetric form is applicable, it is convenient in MINRES and MINRES-QLP; see (3), (5), and (15). We also remark that MINRES and MINRES-QLP produce good estimates of the largest and smallest singular values of $\bar{A}$ (via diagonal values of R_k or L_k in (7) and (11); see [Choi et al. 2011, Section 4]).

Three other regularization tools in the literature (see [Golub and Van Loan 1996, Sections 12.1.1-12.1.3] and [Hansen 1998]) are LSQI, cross-validation, and L-curve. LSQI involves solving a nonlinear equation and is not immediately compatible with the Lanczos framework. Cross-validation takes one row out at a time and thus does not preserve symmetry. The L-curve approach for a CG-type method takes iteration k as the regularization parameter [Hansen 1998, Chapter 8] if both $\|r_k\|$ and $\|x_k\|$ are monotonic. By design, $\|r_k\|$ is monotonic in MINRES and MINRES-QLP, and so is $\|x_k\|$ when $\bar{A}$ is positive definite [Fong 2011]. Otherwise, we prefer the condition L-curve approach in [Calvetti et al. 2000], which graphs cond(T_k) against $\|r_k\|$. Yet another L-curve feasible in MINRES-QLP is $\|x_{k-2}^{(2)}\|$ against $\|r_k\|$, since the former is also monotonic (but available two iterations in lag); see Section 2.4.

2.   MATHEMATICAL BACKGROUND

Notation and details of algorithmic development from [Choi 2006; Choi et al. 2011] are summarized here. As noted earlier, "A" in (1)–(3) is treated as $A - \sigma I$.

2.1   Lanczos Process

MINRES and MINRES-QLP use the symmetric Lanczos process [Lanczos 1950] to reduce A to a tridiagonal form $\underline{T_k}$. The process is initialized with $v_0 \equiv 0$, $\beta_1 = \|b\|$, and $\beta_1 v_1 = b$. After k steps of the tridiagonalization, we have produced

$$p_k = Av_k - \sigma v_k, \qquad \alpha_k = v_k^T p_k, \qquad \beta_{k+1}v_{k+1} = p_k - \alpha_k v_k - \beta_k v_{k-1}, \quad (5)$$

where we choose $\beta_k > 0$ to give $\|v_k\| = 1$. Numerically,

$$p_k = Av_k - \sigma v_k - \beta_k v_{k-1}, \qquad \alpha_k = v_k^T p_k, \qquad \beta_{k+1}v_{k+1} = p_k - \alpha_k v_k$$

is slightly better than (5) [Paige 1976], but we can express (5) in matrix form:

$$V_k \equiv \begin{bmatrix} v_1 \cdots v_k \end{bmatrix}, \qquad AV_k = V_{k+1}\underline{T_k}, \qquad \underline{T_k} \equiv \begin{bmatrix} T_k \\ \beta_{k+1}e_k^T \end{bmatrix}, \qquad (6)$$

where $T_k = \mathrm{tridiag}(\beta_i, \alpha_i, \beta_{i+1})$, $i = 1, \ldots, k$. In exact arithmetic, the Lanczos vectors in the columns of V_k are orthonormal, and the process stops with $k = \ell$ when $\beta_{\ell+1} = 0$ for some $\ell \leq n$, and then $AV_\ell = V_\ell T_\ell$. The rank of T_ℓ could be ℓ or $\ell - 1$ (see Theorem 2.2).

2.2   MINRES Phase

MINRES-QLP typically starts with a MINRES phase, which applies a series of reflectors Q_k to transform $\underline{T_k}$ to an upper triangular matrix $\underline{R_k}$:

$$Q_k \begin{bmatrix} \underline{T_k} & \beta_1 e_1 \end{bmatrix} = \begin{bmatrix} R_k & t_k \\ 0 & \phi_k \end{bmatrix} \equiv \begin{bmatrix} \underline{R_k} & \bar{t}_{k+1} \end{bmatrix}, \qquad (7)$$

where

$$Q_k = Q_{k,k+1}\begin{bmatrix} Q_{k-1} & \\ & 1 \end{bmatrix}, \qquad Q_{k,k+1} \equiv \begin{bmatrix} I_{k-1} & & \\ & c_k & s_k \\ & s_k & -c_k \end{bmatrix}.$$

In the kth step, $Q_{k,k+1}$ is effectively a Householder reflector of dimension 2 [Trefethen and Bau 1997, Exercise 10.4], and its action including its effect on later columns of T_j, $k < j \leq \ell$, is compactly described by

$$\begin{bmatrix} c_k & s_k \\ s_k & -c_k \end{bmatrix}\begin{bmatrix} \gamma_k & \delta_{k+1} & 0 & \Big| & \phi_{k-1} \\ \beta_{k+1} & \alpha_{k+1} & \beta_{k+2} & \Big| & 0 \end{bmatrix} = \begin{bmatrix} \gamma_k^{(2)} & \delta_{k+1}^{(2)} & \epsilon_{k+2} & \Big| & \tau_k \\ 0 & \gamma_{k+1} & \delta_{k+2} & \Big| & \phi_k \end{bmatrix},$$

where the superscripts with numbers in parentheses indicate the number of times the values have been modified. The kth solution approximation to (3) is then defined to be $x_k = V_k y_k$, where y_k solves the subproblem

$$y_k = \arg\min_{y \in \mathbb{R}^k} \|\underline{T_k}y - \beta_1 e_1\| = \arg\min_{y \in \mathbb{R}^k} \|\underline{R_k}y - \bar{t}_{k+1}\|. \qquad (8)$$

When $k < \ell$, R_k is nonsingular and the unique solution of the above subproblem satisfies $R_k y_k = t_k$. Instead of solving for y_k, MINRES solves $R_k^T D_k^T = V_k^T$ by forward substitution, obtaining the last column d_k of D_k at iteration k. At the same time, it updates $x_k \in \mathcal{K}_k(A, b)$ (see Table I for definition) via $x_0 \equiv 0$ and

$$x_k = V_k y_k = D_k R_k y_k = D_k t_k = x_{k-1} + \tau_k d_k, \quad \tau_k \equiv e_k^T t_k, \qquad (9)$$

where one can show using $V_k = D_k R_k$ that $d_k = (v_k - \delta_k^{(2)} d_{k-1} - \epsilon_k d_{k-2})/\gamma_k^{(2)}$.

2.3 MINRES-QLP Phase

The MINRES phase transfers to the MINRES-QLP phase when an estimate of the condition number of A exceeds an input parameter *trancond*. Thus, $trancond > 1/\varepsilon$ leads to MINRES iterates throughout (where $\varepsilon \approx 10^{-16}$ denotes the floating-point precision), whereas $trancond = 1$ generates MINRES-QLP iterates from the start.

Suppose for now that there is no MINRES phase. Then MINRES-QLP applies left reflections as in (7) and a further series of right reflections to transform R_k to a lower triangular matrix $L_k = R_k P_k$, where

$$P_k = P_{1,2} \ \ P_{1,3} P_{2,3} \ \cdots \ \ P_{k-2,k} P_{k-1,k},$$

$$P_{k-2,k} = \begin{bmatrix} I_{k-3} & & \\ & c_{k2} & & s_{k2} \\ & & 1 & \\ & s_{k2} & & -c_{k2} \end{bmatrix}, \qquad P_{k-1,k} = \begin{bmatrix} I_{k-2} & & \\ & c_{k3} & s_{k3} \\ & s_{k3} & -c_{k3} \end{bmatrix}.$$

In the kth step, the actions of $P_{k-2,k}$ and $P_{k-1,k}$ are compactly described by

$$
\begin{bmatrix} \gamma_{k-2}^{(5)} & & \epsilon_k \\ \vartheta_{k-1} & \gamma_{k-1}^{(4)} & \delta_k^{(2)} \\ & & \gamma_k^{(2)} \end{bmatrix}
\begin{bmatrix} c_{k2} & & s_{k2} \\ & 1 & \\ s_{k2} & & -c_{k2} \end{bmatrix}
\begin{bmatrix} 1 & & \\ & c_{k3} & s_{k3} \\ & s_{k3} & -c_{k3} \end{bmatrix}
$$

$$
= \begin{bmatrix} \gamma_{k-2}^{(6)} & & \\ \vartheta_{k-1}^{(2)} & \gamma_{k-1}^{(4)} & \delta_k^{(3)} \\ \eta_k & & \gamma_k^{(3)} \end{bmatrix}
\begin{bmatrix} 1 & & \\ & c_{k3} & s_{k3} \\ & s_{k3} & -c_{k3} \end{bmatrix}
= \begin{bmatrix} \gamma_{k-2}^{(6)} & & \\ \vartheta_{k-1}^{(2)} & \gamma_{k-1}^{(5)} & \\ \eta_k & \vartheta_k & \gamma_k^{(4)} \end{bmatrix}. \tag{10}
$$

The kth approximate solution to (3) is then defined to be $x_k = V_k y_k = V_k P_k u_k = W_k u_k$, where u_k solves the subproblem

$$u_k \equiv \arg\min_u \|u\| \quad \text{s.t.} \quad u \in \arg\min_{u \in \mathbb{R}^k} \left\| \begin{bmatrix} L_k \\ 0 \end{bmatrix} u - \begin{bmatrix} t_k \\ \phi_k \end{bmatrix} \right\|. \tag{11}$$

For $k < \ell$, R_k and L_k are nonsingular because $\underline{T_k}$ has full column rank by Lemma 2.1 below. It is only when $k = \ell$ and $b \notin \text{range}(A)$ that R_k and L_k are singular with rank $\ell - 1$ by Theorem 2.2, in which case one can show that $\eta_k = \gamma_k^{(3)} = \vartheta_k = \gamma_k^{(4)} = 0$ in (10) and $L_\ell = \begin{bmatrix} L_{\ell-1} & 0 \\ 0 & 0 \end{bmatrix}$ with $L_{\ell-1}$ nonsingular. In any case, we need to solve only the nonsingular lower triangular systems $L_k u_k = t_k$ or $L_{\ell-1} u_{\ell-1} = t_{\ell-1}$. Then, u_k and $y_k = P_k u_k$ are the min-length solutions of (11) and (8), respectively.

MINRES-QLP updates x_{k-2} to obtain x_k by short-recurrence orthogonal steps:

$$x_{k-2}^{(2)} = x_{k-3}^{(2)} + \mu_{k-2}^{(3)} w_{k-2}^{(4)}, \text{ where } x_{k-3}^{(2)} \equiv W_{k-3}^{(4)} u_{k-3}^{(3)}, \tag{12}$$

$$x_k = x_{k-2}^{(2)} + \mu_{k-1}^{(2)} w_{k-1}^{(3)} + \mu_k w_k^{(2)}, \tag{13}$$

where w_j refers to the jth column of $W_k = V_k P_k$ and μ_i is the ith element of u_k.

If this phase is preceded by a MINRES phase of k iterations ($0 < k < \ell$), it starts by transferring the last three vectors d_{k-2}, d_{k-1}, d_k to w_{k-2}, w_{k-1}, w_k, and the solution estimate x_k from (9) to $x_{k-2}^{(2)}$ in (12). This needs the last two rows of $L_k u_k = t_k$ (to give μ_{k-1}, μ_k) and the relations $W_k = D_k L_k$ and $x_{k-2}^{(2)} = x_k - \mu_{k-1} w_{k-1} - \mu_k w_k$. The cheaply available right reflections P_k and the bottom right 3×3 submatrix of L_k (i.e., the last term in (10)) need to have been saved in the MINRES phase in order to facilitate the transfer.

2.4 Norm Estimates and Stopping Conditions

Short-term recurrences are used to estimate the following quantities (where we assume $\sigma = 0$ for simplicity):

$$\|r_k\| \approx \phi_k = \phi_{k-1}s_k, \qquad\qquad \phi_0 = \|b\| \qquad\qquad (\phi_k \searrow)$$

$$\|Ar_k\| \approx \psi_k = \phi_k\|[\gamma_{k+1}\ \ \delta_{k+2}]\|, \qquad\qquad (\psi_\ell = 0)$$

$$\|x_k^{(2)}\| \approx \chi_{k-2}^{(2)} = \|[\chi_{k-3}^{(2)}\ \ \mu_{k-2}^{(3)}]\|, \qquad \chi_{-2} = \chi_{-1} = 0 \qquad (\chi_{k-2}^{(2)} \nearrow)$$

$$\|x_k\| \approx \chi_k = \|[\chi_{k-2}^{(2)}\ \ \mu_{k-1}^{(2)}\ \ \mu_k]\|, \qquad \chi_0 = 0 \qquad (\chi_\ell = \|x^\dagger\|)$$

$$\|Ax_k\| \approx \omega_k = \|[\omega_{k-1}\ \ \tau_k]\|, \qquad\qquad \omega_0 = 0 \qquad (\omega_k \nearrow)$$

$$\|A\| \approx \mathcal{A}_k = \max\left\{\mathcal{A}_{k-1}, \|\underline{T_k}e_k\|, \overline{\gamma}_k\right\}, \qquad \mathcal{A}_0 = 0 \qquad (\mathcal{A}_k \nearrow \|A\|)$$

$$\mathrm{cond}(A) \approx \kappa_k = \mathcal{A}_k/\underline{\gamma}_k, \qquad\qquad \kappa_0 = 1 \qquad (\kappa_k \nearrow \mathrm{cond}(A))$$

where $\overline{\gamma}_k$ and $\underline{\gamma}_k$ are the largest and smallest diagonals of L_k in absolute value. The up (down) arrows in parentheses indicate that the associated quantities increase (decrease) monotonically. The last two estimates tend to their targets from below; see [Choi 2006; Choi et al. 2011] for derivation.

MINRES-QLP has 14 possible stopping conditions in five classes that use the above estimates and optional user-input parameters *itnlim*, *rtol*, *Acondlim*, and *maxxnorm*:

(C1) From Lanczos and the QLP factorization:

$$k = \textit{itnlim}; \qquad \beta_{k+1} < \varepsilon; \qquad \left|\gamma_k^{(4)}\right| < \varepsilon;$$

(C2) Normwise relative backward errors (NRBE) [Paige and Strakoš 2002]:

$$\|r_k\|/\left(\|A\|\|x_k\| + \|b\|\right) \le \max(\textit{rtol}, \varepsilon); \qquad \|Ar_k\|/\left(\|A\|\|r_k\|\right) \le \max(\textit{rtol}, \varepsilon);$$

(C3) Regularization attempts:

$$\mathrm{cond}(A) \ge \min(\textit{Acondlim}, \textit{0.1}/\varepsilon); \qquad \|x_k\| \ge \textit{maxxnorm};$$

(C4) Degenerate cases:

$$\beta_1 = 0 \qquad \Rightarrow \qquad b = 0 \qquad \Rightarrow \qquad x = 0 \text{ is the solution;}$$
$$\beta_2 = 0 \qquad \Rightarrow \qquad v_2 = 0 \qquad \Rightarrow \qquad Ab = \alpha_1 b,$$

   i.e., b and α_1 are an eigenpair of A, and $x = b/\alpha_1$ solves $Ax = b$;

(C5) Erroneous inputs:

   A not symmetric; M not symmetric; M not positive definite;

where M is a preconditioner to be described in the next section. For symmetry of A, it is not practical to check $e_i^T A e_j = e_j^T A e_i$ for all $i, j = 1, \ldots, n$. Instead, we statistically test whether $z = |x^T(Ay) - y^T(Ax)|$ is sufficiently small for two nonzero n-vectors x and y (e.g., each element in the vectors is drawn from the standard normal distribution). For positive definiteness of M, since M is positive definite if and only if M^{-1} is positive definite, we simply test that $z_k^T M^{-1} z_k = z_k^T q_k > 0$ each iteration (see Section 3).

We find that the recurrence relations for ϕ_k and ψ_k hold to high accuracy. Thus x_k is an acceptable solution of (3) if the computed value of ϕ_k or ψ_k is suitably small according to the NRBE tests in class (C2) above. When a condition in (C3) is met, the final x_k may or may not be an acceptable solution.

The class (C1) tests for small β_{k+1} and $\gamma_k^{(4)}$ are included in the unlikely case in practice that the theoretical Lanczos termination occurs. Ideally one of the NRBE tests should cause MINRES-QLP to terminate. If not, it is an indication that the problem is very ill-conditioned, in which case the regularization and preconditioning techniques of Sections 1.2 and 3 may be helpful.

2.5 Two Theorems

We complete this section by presenting two theorems from [Choi et al. 2011] with slightly simpler proofs.

LEMMA 2.1. $\mathrm{rank}(\underline{T_k}) = k$ for all $1 \leq k < \ell$.

PROOF. For $1 \leq k < \ell$ we have $\beta_2, \ldots, \beta_{k+1} > 0$ by definition. Hence $\underline{T_k}$ has full column rank. $\square$

THEOREM 2.2. T_ℓ is nonsingular if and only if $b \in \mathrm{range}(A)$. Furthermore, $\mathrm{rank}(T_\ell) = \ell - 1$ if $b \notin \mathrm{range}(A)$.

PROOF. We use $AV_\ell = V_\ell T_\ell$ twice. First, if T_ℓ is nonsingular, we can solve $T_\ell y_\ell = \beta_1 e_1$ and then $AV_\ell y_\ell = V_\ell T_\ell y_\ell = V_\ell \beta_1 e_1 = b$. Conversely, if $b \in \mathrm{range}(A)$, then $\mathrm{range}(V_\ell) \subseteq \mathrm{range}(A)$. Suppose T_ℓ is singular. Then there exists $z \neq 0$ such that $V_\ell T_\ell z = AV_\ell z = 0$. That is, $0 \neq V_\ell z \in \mathrm{null}(A)$. But this is impossible because $V_\ell z \in \mathrm{range}(A)$ and $\mathrm{null}(A) \cap \mathrm{range}(V_\ell) = 0$. Thus, T_ℓ must be nonsingular.

We have shown that if $b \notin \mathrm{range}(A)$, $T_\ell = \begin{bmatrix} \underline{T_{\ell-1}} & \beta_\ell e_{\ell-1} \\ & \alpha_\ell \end{bmatrix}$ is singular, and therefore $\ell > \mathrm{rank}(T_\ell) \geq \mathrm{rank}(\underline{T_{\ell-1}}) = \ell - 1$ by Lemma 2.1. Therefore, $\mathrm{rank}(T_\ell) = \ell - 1$. $\square$

By Lemma 2.1 and Theorem 2.2 we are assured that the QLP decomposition without column pivoting [Stewart 1999; Choi et al. 2011] for $\underline{T_k}$ is rank-revealing, which is a necessary precondition for solving a least-squares problem.

THEOREM 2.3. In MINRES-QLP, x_ℓ is the minimum-length solution of (3).

PROOF. y_ℓ comes from the min-length LS solution of $T_\ell y_\ell \approx \beta_1 e_1$ and thus satisfies the normal equation $T_\ell^2 y_\ell = T_\ell \beta_1 e_1$ and $y_\ell \in \mathrm{range}(T_\ell)$. Now $x_\ell = V_\ell y_\ell$ and $Ax_\ell = AV_\ell y_\ell = V_\ell T_\ell y_\ell$. Hence $A^2 x_\ell = AV_\ell T_\ell y_\ell = V_\ell T_\ell^2 y_\ell = V_\ell T_\ell \beta_1 e_1 = Ab$. Thus x_ℓ is an LS solution of (3). Since $y_\ell \in \mathrm{range}(T_\ell)$, $y_\ell = T_\ell z$ for some z, and so $x_\ell = V_\ell y_\ell = V_\ell T_\ell z = AV_\ell z \in \mathrm{range}(A)$ is the min-length LS solution of (3). $\square$

3. PRECONDITIONING

Iterative methods can be accelerated if preconditioners are available and well-chosen. For MINRES-QLP, we want to choose a symmetric positive-definite matrix M to solve a nonsingular system (1) by implicitly solving an equivalent symmetric consistent system $M^{-\frac{1}{2}} A M^{-\frac{1}{2}} \bar{x} = \bar{b}$, where $M^{\frac{1}{2}} x = \bar{x}$, $\bar{b} = M^{-\frac{1}{2}} b$, and $\mathrm{cond}(M^{-\frac{1}{2}} A M^{-\frac{1}{2}}) \ll \mathrm{cond}(A)$. This two-sided preconditioning preserves symmetry. Thus we can derive preconditioned MINRES-QLP by applying MINRES-QLP to the equivalent problem and setting $x = M^{-\frac{1}{2}} \bar{x}$.

With preconditioned MINRES-QLP, we can solve a singular consistent system (2), but we will obtain a least-squares solution that is not necessarily the minimum-length solution (unless $M = I$). For inconsistent systems (3), preconditioning alters the least-squares norm to $\| \cdot \|_{M^{-1}}$, and the solution is of minimum length in the new norm space. We refer readers to [Choi et al. 2011, Section 7] for a detailed discussion of various approaches to preserving the two-norm "minimum length."

To derive MINRES-QLP, we define

$$z_k = \beta_k M^{\frac{1}{2}} v_k, \qquad q_k = \beta_k M^{-\frac{1}{2}} v_k, \qquad \text{so that} \quad Mq_k = z_k. \qquad (14)$$

Then $\beta_k = \|\beta_k v_k\| = \|M^{-\frac{1}{2}} z_k\| = \|z_k\|_{M^{-1}} = \|q_k\|_M = \sqrt{q_k^T z_k}$, where the square root is well defined because M is positive definite, and the following expressions replace the quantities in (5) in the Lanczos iterations:

$$p_k = Aq_k - \sigma q_k, \qquad \alpha_k = \frac{1}{\beta_k^2} q_k^T p_k, \qquad z_{k+1} = \frac{1}{\beta_k} p_k - \frac{\alpha_k}{\beta_k} z_k - \frac{\beta_k}{\beta_{k-1}} z_{k-1}. \quad (15)$$

We also need to solve the system $Mq_k = z_k$ in (14) at each iteration.

In the MINRES phase, we define $\bar{d}_k = M^{-\frac{1}{2}} d_k$ and update the solution of the original problem (1) by

$$\bar{d}_k = \left(\frac{1}{\beta_k} q_k - \delta_k^{(2)} \bar{d}_{k-1} - \epsilon_k \bar{d}_{k-2} \right) \big/ \gamma_k^{(2)}, \qquad x_k = M^{-\frac{1}{2}} \bar{x}_k = x_{k-1} + \tau_k \bar{d}_k.$$

In the MINRES-QLP phase, we define $\overline{W}_k \equiv M^{-\frac{1}{2}} W_k = (M^{-\frac{1}{2}} V_k) P_k$ and update the solution estimate of problem (1) by orthogonal steps:

$$\begin{aligned} \bar{w}_k &= -(c_{k2}/\beta_k)q_k + s_{k2}\bar{w}_{k-2}^{(3)}, & \bar{w}_{k-2}^{(4)} &= (s_{k2}/\beta_k)q_k + c_{k2}\bar{w}_{k-2}^{(3)}, \\ \bar{w}_k^{(2)} &= s_{k3}\bar{w}_{k-1}^{(2)} - c_{k3}\bar{w}_k, & \bar{w}_{k-1}^{(3)} &= c_{k3}\bar{w}_{k-1}^{(2)} + s_{k3}\bar{w}_k, \\ x_{k-2}^{(2)} &= x_{k-3}^{(2)} + \mu_{k-2}^{(3)}\bar{w}_{k-2}^{(4)}, & x_k &= x_{k-2}^{(2)} + \mu_{k-1}^{(2)}\bar{w}_{k-1}^{(3)} + \mu_k\bar{w}_k^{(2)}. \end{aligned}$$

Let $\bar{r}_k = \bar{b} - M^{-\frac{1}{2}}(A - \sigma I)M^{-\frac{1}{2}}\bar{x}_k = M^{-\frac{1}{2}} r_k$. Then $x_k = M^{-\frac{1}{2}}\bar{x}_k$ is an acceptable solution of (1) if the computed value of $\phi_k \approx \|\bar{r}_k\| = \|r_k\|_{M^{-1}}$ is sufficiently small.

We can now present our pseudocode in Algorithm 1. The reflectors are implemented in Algorithm 2 SymOrtho(a, b) for real a and b, which is a stable form for computing $r = \sqrt{a^2 + b^2} \geq 0$, $c = \frac{a}{r}$, and $s = \frac{b}{r}$. The complexity is at most 6 flops and a square root. Algorithm 1 lists all steps of MINRES-QLP with preconditioning. For simplicity, $\bar{w}_k$ is written as w_k for all relevant k. Also, the output x solves $(A - \sigma I)x \approx b$, but other outputs are associated with the preconditioned system.

4. KEY FORTRAN 90 DESIGN FEATURES

In this section we describe the key features of our FORTRAN implementations. For an accessible reference on the language syntax in various FORTRAN releases, we refer readers to [Chivers and Sleightholme 2006].

Our FORTRAN 90 package contains the following files for symmetric problems with the first three files forming the core. Their dependencies are depicted in Figure 1.

1. `minresqlpDataModule.f90`: defines precision and constants used in other modules; see Listing 1

Algorithm 1: Pseudocode of preconditioned MINRES-QLP for solving $(A - \sigma I)x \approx b$. In the right-justified comments, $\tilde{A} \equiv M^{-\frac{1}{2}}(A - \sigma I)M^{-\frac{1}{2}}$.

---

input: A, b, σ, M

1   $z_0 = 0, \quad z_1 = b, \quad$ Solve $Mq_1 = z_1, \quad \beta_1 = \sqrt{b^T q_1}, \quad \phi_0 = \beta_1$ [Initialize]

2   $w_0 = w_{-1} = 0, \quad x_{-2} = x_{-1} = x_0 = 0$

3   $c_{0,1} = c_{0,2} = c_{0,3} = -1, \quad s_{0,1} = s_{0,2} = s_{0,3} = 0, \quad \tau_0 = \omega_0 = \chi_{-2} = \chi_{-1} = \chi_0 = 0$

4   $\kappa_0 = 1, \quad \mathcal{A}_0 = \delta_1 = \gamma_{-1} = \gamma_0 = \eta_{-1} = \eta_0 = \eta_1 = \vartheta_{-1} = \vartheta_0 = \vartheta_1 = \mu_{-1} = \mu_0 = 0$

5   $k = 0$

6   **while** no stopping condition is satisfied **do**

7     $k \leftarrow k + 1$

8     $p_k = Aq_k - \sigma q_k, \qquad \alpha_k = \frac{1}{\beta_k^2} q_k^T p_k$ [Preconditioned Lanczos]

9     $z_{k+1} = \frac{1}{\beta_k} p_k - \frac{\alpha_k}{\beta_k} z_k - \frac{\beta_k}{\beta_{k-1}} z_{k-1}$

10   Solve $Mq_{k+1} = z_{k+1}, \qquad \beta_{k+1} = \sqrt{q_{k+1}^T z_{k+1}}$

11   **if** $k = 1$ **then** $\rho_k = \|[\alpha_k \ \ \beta_{k+1}]\|$ **else** $\rho_k = \|[\beta_k \ \ \alpha_k \ \ \beta_{k+1}]\|$

12    $\delta_k^{(2)} = c_{k-1,1}\delta_k + s_{k-1,1}\alpha_k$ [Previous left reflection...]

13    $\gamma_k = s_{k-1,1}\delta_k - c_{k-1,1}\alpha_k$ [on middle two entries of $\underline{T_k}e_k$...]

14    $\epsilon_{k+1} = s_{k-1,1}\beta_{k+1}$ [produces first two entries in $\underline{T_{k+1}}e_{k+1}$]

15    $\delta_{k+1} = -c_{k-1,1}\beta_{k+1}$

16    $c_{k1}, s_{k1}, \gamma_k^{(2)} \leftarrow \text{SymOrtho}(\gamma_k, \beta_{k+1})$ [Current left reflection]

17    $c_{k2}, s_{k2}, \gamma_{k-2}^{(6)} \leftarrow \text{SymOrtho}(\gamma_{k-2}^{(5)}, \epsilon_k)$ [First right reflection]

18    $\delta_k^{(3)} = s_{k2}\vartheta_{k-1} - c_{k2}\delta_k^{(2)}, \qquad \gamma_k^{(3)} = -c_{k2}\gamma_k^{(2)}, \qquad \eta_k = s_{k2}\gamma_k^{(2)}$

19    $\vartheta_{k-1}^{(2)} = c_{k2}\vartheta_{k-1} + s_{k2}\delta_k^{(2)}$

20    $c_{k3}, s_{k3}, \gamma_{k-1}^{(5)} \leftarrow \text{SymOrtho}(\gamma_{k-1}^{(4)}, \delta_k^{(3)})$ [Second right reflection...]

21    $\vartheta_k = s_{k3}\gamma_k^{(3)}, \qquad \gamma_k^{(4)} = -c_{k3}\gamma_k^{(3)}$ [to zero out $\delta_k^{(3)}$]

22    $\tau_k = c_{k1}\phi_{k-1}$ [Last element of t_k]

23    $\phi_k = s_{k1}\phi_{k-1}, \quad \psi_{k-1} = \phi_{k-1}\|[\gamma_k \ \ \delta_{k+1}]\|$ [Update $\|\bar{r}_k\|$, $\|\tilde{A}\bar{r}_{k-1}\|$]

24   **if** $k = 1$ **then** $\gamma_{\min} = \gamma_1$ **else** $\gamma_{\min} \leftarrow \min\{\gamma_{\min}, \gamma_{k-2}^{(6)}, \gamma_{k-1}^{(5)}, |\gamma_k^{(4)}|\}$

25    $\mathcal{A}_k = \max\{\mathcal{A}_{k-1}, \rho_k, \gamma_{k-2}^{(6)}, \gamma_{k-1}^{(5)}, |\gamma_k^{(4)}|\}$ [Update $\|\tilde{A}\|$]

26    $\omega_k = \|[\omega_{k-1} \ \ \tau_k]\|, \qquad \kappa_k \leftarrow \mathcal{A}_k/\gamma_{\min}$ [Update $\|\tilde{A}x_k\|$, $\text{cond}(\tilde{A})$]

27    $w_k = -(c_{k2}/\beta_k)q_k + s_{k2}w_{k-2}^{(3)}$ [Update w_{k-2}, w_{k-1}, w_k]

28    $w_{k-2}^{(4)} = (s_{k2}/\beta_k)q_k + c_{k2}w_{k-2}^{(3)}$

29   **if** $k > 2$ **then** $w_k^{(2)} = s_{k3}w_{k-1}^{(2)} - c_{k3}w_k, \qquad w_{k-1}^{(3)} = c_{k3}w_{k-1}^{(2)} + s_{k3}w_k$

30   **if** $k > 2$ **then** $\mu_{k-2}^{(3)} = (\tau_{k-2} - \eta_{k-2}\mu_{k-4}^{(4)} - \vartheta_{k-2}\mu_{k-3}^{(3)})/\gamma_{k-2}^{(6)}$ [Update μ_{k-2}]

31   **if** $k > 1$ **then** $\mu_{k-1}^{(2)} = (\tau_{k-1} - \eta_{k-1}\mu_{k-3}^{(3)} - \vartheta_{k-1}^{(2)}\mu_{k-2}^{(3)})/\gamma_{k-1}^{(5)}$ [Update μ_{k-1}]

32   **if** $\gamma_k^{(4)} \neq 0$ **then** $\mu_k = (\tau_k - \eta_k\mu_{k-2}^{(3)} - \vartheta_k\mu_{k-1}^{(2)})/\gamma_k^{(4)}$ **else** $\mu_k = 0$ [Compute μ_k]

33    $x_{k-2}^{(2)} = x_{k-3}^{(2)} + \mu_{k-2}^{(3)}w_{k-2}^{(3)}$ [Update x_{k-2}]

34    $x_k = x_{k-2}^{(2)} + \mu_{k-1}^{(2)}w_{k-1}^{(3)} + \mu_k w_k^{(2)}$ [Compute x_k]

35    $\chi_{k-2}^{(2)} = \|[\chi_{k-3}^{(2)} \ \ \mu_{k-2}^{(3)}]\|$ [Update $\|x_{k-2}\|$]

36    $\chi_k = \|[\chi_{k-2}^{(2)} \ \ \mu_{k-1}^{(2)} \ \ \mu_k]\|$ [Compute $\|x_k\|$]

37   $x = x_k, \quad \phi = \phi_k, \quad \psi = \phi_k\|[\gamma_{k+1} \ \ \delta_{k+2}]\|, \quad \chi = \chi_k, \quad \mathcal{A} = \mathcal{A}_k, \quad \omega = \omega_k$

    output: $x, \phi, \psi, \chi, \mathcal{A}, \kappa, \omega$

Algorithm 2: Algorithm SymOrtho.

---

input: a, b

1 **if** $b = 0$ **then** $s = 0$, $\qquad r = |a|$
2 $\qquad$ **if** $a = 0$ **then** $c = 1$ **else** $c = \text{sign}(a)$
3 **else if** $a = 0$ **then**
4 $\quad\big\lfloor\ c = 0$, $\qquad s = \text{sign}(b)$, $\qquad r = |b|$
5 **else if** $|b| \geq |a|$ **then**
6 $\quad\big\lfloor\ \tau = a/b$, $\qquad s = \text{sign}(b)/\sqrt{1 + \tau^2}$, $\qquad c = s\tau$, $\qquad r = b/s$
7 **else if** $|a| > |b|$ **then**
8 $\quad\big\lfloor\ \tau = b/a$, $\qquad c = \text{sign}(a)/\sqrt{1 + \tau^2}$, $\qquad s = c\tau$, $\qquad r = a/c$
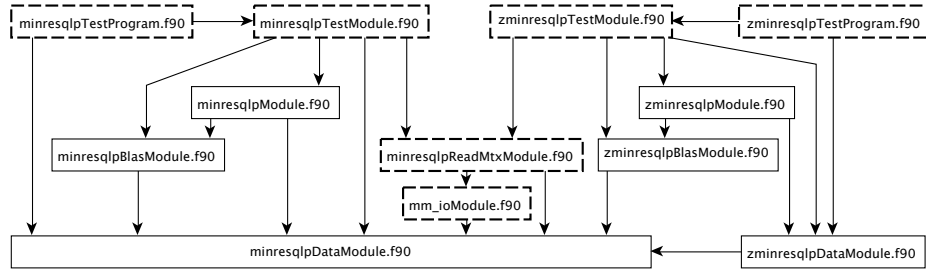
output: c, s, r

Fig. 1. FORTRAN 90 source files and their dependencies. Filenames boxed in broken lines are optional, and the corresponding files are used mainly for testing and demonstration.

2. `mminresqlpBlasModule.f90`: packages some BLAS functions [Burkardt]

3. `minresqlpModule.f90`: implements MINRES-QLP with preconditioning; see Listing 2 for partial source code

4. `mm_ioModule.f90` and `minresqlpReadMtxModule.f90`: packages subroutines for reading Matrix Market files [Burkardt]

5. `minresqlpTestModule.f90`: illustrates how MINRES-QLP can call `Aprod` or `Msolve` with a short fixed parameter list, even if it needs arbitrary other data; see Listing 3 for partial source code

6. `minresqlpTestProgram.f90`: contains the main driver program of unit tests

7. `Makefile`: compiles the FORTRAN source files via the Unix command `make`

8. `readme.txt`: contains information about software license, other files in the package, and program compilation and execution.

The counterparts of these programs for Hermitian problems have the same filenames prefixed with the letter "`z`".

We review and step through the code in the following subsections. The line numbers in Listings 1–3 are used for reference only and do not correspond to actual line numbers in the source code. The vertical dots in Listing 2 lines 35 and 43 indicate omitted code of one or more lines. We also note that FORTRAN 90 keywords are displayed in bold in the listings, and that comments are marked with exclamation marks in italics.

Listing 1.   Partial FORTRAN 90 code listing of minresqlpDataModule.

```
1  module minresqlpDataModule
2   implicit none

4   intrinsic                         ::        selected_real_kind

6   integer,         parameter, public :: dp = selected_real_kind(15)

7   real(kind=dp), parameter, public :: zero = 0.0_dp, one = 1.0_dp

8  end module minresqlpDataModule
```

4.1 Overloaded Intrinsic Operators and BLAS Procedures

For standard vector operations, we simply apply the intrinsic arithmetic and assignment operators $\pm, \times, =$. In addition we adopt a FORTRAN 90 translation [Burkardt] of two external level-1 BLAS functions ddot and dnrm2 [BLAS] for computing inner products and two norms of vectors, which take care to avoid undesirable overflow or underflow.

In minresqlpModule, the line "use minresqlpBlasModule" can be omitted if the code is already linked to a BLAS library.

4.2 Using Modules and Interface and Passing User-Defined Subroutines to MINRESQLP

In our FORTRAN 90 implementation, we use *modules* instead of the obsolete FORTRAN 77 COMMON blocks for grouping programs units and data together and controlling their availability to other program units. A module can use public data and subroutines from other modules (by declaring an *interface* block), share its own public data and subroutines with other program units, and hide its own private data and subroutines from being used by other program units. We can also use modules to package procedures.

In Listing 2, line 2, module minresqlpModule uses the external public constant dp from minresqlpDataModule. From line 9 onwards, minresqlpModule defines a public subroutine MINRESQLP, where we implement MINRES-QLP in Algorithm 1.

A FORTRAN subroutine may have multiple and optional input and output arguments, which transfer information to and from a calling program. MINRESQLP has a total of 20 arguments (see lines 9-12). The data types and intent of these arguments are declared in lines 15-25. For example, the first argument n in line 15 is an input integer, whereas x(n) in line 23 is an output n-vector of double precision.

Two input arguments Aprod and Msolve are external user-defined subroutines (Listing 3, lines 8 and 18-32) being passed into MINRESQLP as inputs—we recommend they be private for data integrity. The subroutine Aprod defines the matrix A as an operator (in Algorithm 1, line 8). For a given vector x, the FORTRAN statement call Aprod(n, x, y) must return the product $y = Ax$ without altering the vector x. The subroutine Msolve is optional, and it defines a symmetric positive-definite matrix as an operator M that serves as a preconditioner (line 10 in Algorithm 1). For a given vector y, the FORTRAN statement call MSolve(n,

Listing 2.   Partial code listing of subroutine MINRESQLP in minresqlpModule.

```fortran
1  module minresqlpModule
2    use  minresqlpDataModule , only : dp , one , zero
3    use  minresqlpBlasModule , only : dnrm2 , ddot

5    implicit none

7    public    :: MINRESQLP, SYMORTHO
8  contains
9      subroutine MINRESQLP(                                          &
10               n, Aprod , b , shift , Msolve , disable , nout ,      &
11               itnlim , rtol , maxxnorm , trancond , Acondlim ,      &
12               x , istop , itn , rnorm , Arnorm , xnorm , Anorm , Acond )

14      ! Inputs
15      integer(ip) , intent(in)                 :: n
16      real(dp) ,       intent(in)              :: b(n)
17      integer(ip) , intent(in) , optional      :: itnlim , nout
18      logical ,         intent(in) , optional  :: disable
19      real(dp) ,       intent(in) , optional   :: shift
20      real(dp) ,       intent(in) , optional   :: rtol , maxxnorm ,
           trancond , Acondlim

22      ! Outputs
23      real(dp) ,       intent(out)             :: x(n)
24      integer(ip) , intent(out) , optional     :: istop , itn
25      real(dp) ,       intent(out) , optional  :: rnorm , Arnorm , xnorm ,
           Anorm , Acond

28      interface
29       subroutine Aprod (n ,x ,y )                          ! y := Ax
30         use          minresqlpDataModule
31         integer ,  intent(in)      :: n
32         real(dp) , intent(in)      :: x(n)
33         real(dp) , intent(out)     :: y(n)
34       end subroutine Aprod
35       ⋮
36      end interface

38      intrinsic :: abs , epsilon , sqrt

40      ! Local arrays and variables
41      real(dp)  :: r1(n) , r2(n) , v(n) , w(n) , wl(n) ,          &

43      ⋮
44   end subroutine MINRESQLP
45   ⋮
46  end module minresqlpModule
```

Listing 3.   Partial code listing of minresqlpTestModule.

```
 1  module minresqlpTestModule
 2     use   minresqlpDataModule, only : dp
 3     use   minresqlpModule,      only : MINRESQLP

 5     implicit none

 7     public    :: minresqlptest
 8     private   :: Aprod, Msolve

10     ! DYNAMIC WORKSPACE DEFINED HERE.
11     ! It is allocated in minresqlptest and used by Aprod or Msolve.

13     real(dp), allocatable :: d(:)      !Defines diagonal matrix D.
14     real(dp)              :: Ashift    !Shift diagonal elements of D in Msolve.
15     real(dp)              :: Mpert     !Perturbation to D in Msolve
16                                        !to avoid having an exact preconditioner.
17  contains
18     subroutine Aprod(n,x,y)

20        integer,  intent(in)   :: n
21        real(dp), intent(in)   :: x(n)
22        real(dp), intent(out)  :: y(n)

24        integer  :: i

26        do i = 1, n
27           y(i) = d(i)*x(i)
28        end do

30     end subroutine Aprod

    .
32     .
    .

34     subroutine minresqlptest( n, precon, shift, pertM, nout )
       .
35     .
36      call MINRESQLP(n, Aprod, b, shift, Msolve, disable,    &
37            nout, itnlim, rtol, maxxnorm, trancond, Acondlim, &
38            x, istop, itn, rnorm, Arnorm, xnorm, Anorm, Acond )

    .
40     .
41     end subroutine minresqlptest
42  end module minresqlpTestModule
```

y, x) must solve the linear system $Mx = y$ without altering the vector y. To provide the compiler the necessary information about these `private` subroutines defined in `minresqlpTestModule`, an `interface` block in subroutine `MINRESQLP` is declared (lines 28-36 in Listing 2), which essentially replicates the headers of `Aprod` and `Msolve` in `minresqlpTestModule` (lines 18-32 in Listing 3).

`MINRESQLP` is called by the public routine `minresqlptest` defined in module `minresqlpTestModule` (see lines 7, 34-41 in Listing 3). Since `MINRESQLP` is public (Listing 2, line 7), `minresqlpTestModule` can simply `use` it (Listing 3, line 3). We have not listed details of `minresqlptest`, but it calls `MINRESQLP` with `Aprod` and `Msolve` passed as parameters (Listing 3, line 36).

We note that subroutine arrays and variables such as `r1(n)` in Listing 2, line 41, and `i` in Listing 3, line 24, are by default private and not accessible to other program units. In contrast, module arrays and variables are by default public and accessible to other program units. We have marked `d(:)`, `Ashift`, `Mpert` as `private` in Listing 3, lines 13–15, in order to make them accessible to only the subroutines `minresqlptest`, `Aprod`, and `Msolve` in the containing module but not outside.

To summarize, we have described and provided a pattern that allows MINRES-QLP users to solve different problems by simply editing `minresTestModule` (and possibly the main program `minresTestProgram`, which calls `minresqlptest`). Users do not need to change `MINRESQLP` as long as the header of subroutines `Aprod` and `Msolve` stay the same in `minresTestModule`. If necessary, local arrays or variables such as `d(:)` can be used instead of additional input arguments to define these operators. In this way, users can make the data A and M known to `MINRESQLP` but hidden and thus secure from other programs.

Our design spares users from implementing *reverse communication*, and hence enables the development of iterative methods without *a priori* knowledge of users' problem data A and M (by returning control to the calling program every time `Aprod` or `Msolve` is to be invoked). While reverse communication is widely used in scientific computing with FORTRAN 77, the resulting code usually appears formidable and unrecognizable from the original pseudocode; see [Dongarra et al. 1995] and [Oliveira and Stewart 2006] for two examples of CG and numerical integration coded in FORTRAN 77 and 90, respectively. Our MINRES-QLP implementation achieves the purpose of reverse communication while preserving code readability and thus maintainability. The FORTRAN 90 module structure allows the user's Ax products and $Mx = y$ solves to be implemented outside MINRES-QLP in the same way that MATLAB's function handles operate.

4.3 Unit Testing

Unit testing is an important software development strategy that cannot be overemphasized, especially in the scientific computing communities. Unit testing usually consists of multiple small and fast but specific and illuminating test cases that check whether the code behaves as designed. Software development is incremental, and errors (also known as bugs) are often found over time. Adding new functionalities or fixing errors often breaks the code for some earlier successful test cases. It is therefore critical to expand the test cases and to ensure that all unit tests are executed with expected results every time a key program unit is updated.

In our development of FORTRAN 90 MINRES-QLP, we have created a suite of 117

test cases including singular matrices representative of real-world applications [Foster 2009; Davis and Hu 2011]. The test program outputs results to `MINRESQLP.txt`. If users need to modify subroutine `MINRESQLP`, they can run these test cases and search for the word "appear" in the output file to check whether all tests are reported to be successful. For more sophisticated unit testing frameworks employed in large-scale scientific software development, see [O'Boyle et al. 2008].

4.4 Miscellaneous Issues

The complex program units for Hermitian linear systems and LS problems are similar to the real ones, and thus we will not go into detail. Many variables of type `real(dp)` are changed to `complex(dp)`.

To use a different precision throughout the program units, MINRES-QLP users can simply edit the input argument value of `dp` in `minresqlpDataModule`, line 6.

In the main subroutine `MINRESQLP`, we provide a logical parameter `debug` as a diagnostic tool; when it is true, variable values are printed to the standard output.

5. INPUTS, OUTPUTS, AND NUMERICAL EXAMPLES

Subroutine `MINRESQLP` contains the core implementation of MINRES-QLP and has 12 input parameters documented in the code as well as in Table III. It uses seven local n-vectors and returns a computed solution `x` as one of the eight outputs. Mandatory inputs are n, `Aprod`, and b. All outputs other than x are optional. If an input is optional, MINRES-QLP prescribes a default value. It is well known that careful choice of parameter values is critical in the convergence behavior of iterative solvers. Although the default parameter values in MINRES-QLP work well in most tests, they may need to be fine tuned in some cases by trial and error, solving a series of problems as in iterative regularization, or partial or full reorthogonalization of the Lanczos vectors.

Table III: Input parameters in subroutine MINRES-QLP.

Input	Definition
n	The dimension of the symmetric matrix or operator A.
$b(n)$	The right-hand-side vector b.
`Aprod`	An external subroutine defining the matrix A. For a given vector x, the statement `call Aprod ( n, x, y )` must return the product $y = Ax$ without altering the vector x. An extra call of `Aprod` is used to check if A is symmetric. The program calling MINRES-QLP must declare `Aprod` to be external.
`Msolve`	An optional external subroutine defining a preconditioner M,

Table III: Input parameters in MINRES-QLP (continued).

Input	Definition
	which should approximate $A - shiftI$ in some sense. M must be symmetric positive definite. For a given vector x, the statement `call Msolve( n, x, y )` must solve the linear system $My = x$ without altering the vector x. In general, M should be chosen so that $\tilde{A} \equiv M^{-\frac{1}{2}}\bar{A}M^{-\frac{1}{2}}$ has more clustered eigenvalues. If $\bar{A}$ is positive definite, $\tilde{A}$ would ideally be close to a multiple of I. If $\bar{A}$ is indefinite, $\tilde{A}$ might be close to a multiple of `diag`$(I \quad -I)$.
shift	Should be zero if the system $Ax = b$ is to be solved. Otherwise, it could be an approximation to an eigenvalue of A, such as the Rayleigh quotient $(b^TAb)/(b^Tb)$ corresponding to the vector b. If b is sufficiently like an eigenvector corresponding to an eigenvalue near *shift*, then the computed x may have very large components. When normalized, x may be closer to an eigenvector than b.
nout	A file number. The calling program must open a file for output using for example: `open(nout, file='MINRESQLP.txt', status='unknown')`. If $nout > 0$, a summary of the iterations will be printed on unit *nout*. If *nout* is absent or the file associated with *nout* is not open properly, results will be written to `MINRESQLP_tmp.txt`.
itnlim	An upper limit on the number of iterations. Default to $4n$.
rtol	A user-specified tolerance. MINRES-QLP terminates if it appears that $\|\bar{r}\|$ is smaller than $rtol(\|\bar{A}\|\|\bar{x}\| + \|b\|)$, where

Table III: Input parameters in MINRES-QLP (continued).

Input	Definition
	$\bar{r} = \bar{b} - \bar{A}\bar{x}$, or that $\|\bar{A}\bar{r}\|$ is smaller than $rtol\|\bar{A}\|\|\bar{r}\|$.
	If $shift = 0$ and Msolve is absent, MINRES-QLP terminates if
	$\|r\|$ is smaller than $rtol(\|A\|\|x\| + \|b\|)$, where $r = b - Ax$,
	or if $\|Ar\|$ is smaller than $rtol\|A\|\|r\|$. Default to ε.
maxxnorm	An upper bound on $\|x\|$. Default value is 10^7.
Acondlim	An upper bound on $Acond$, an estimate of cond(A).
	Default value is 10^{15}.
trancond	If $trancond > 1$, a switch is made from MINRES iterations
	to MINRES-QLP iterations when $Acond \geq trancond$.
	If $trancond = 1$, all iterations will be MINRES-QLP iterations.
	If $trancond = acondlim$, all iterations will be conventional
	MINRES iterations (which are slightly cheaper).
	Default value is 10^7.

We use two small examples to illustrate the outputs of MINRESQLP. We refer readers to [Choi 2006, Chapter 4] or [Choi et al. 2011, Section 8] for more significant numerical examples.

Table IV compares the MINRES solution to the MINRES-QLP solution for the small problem $Ax \approx b$, where $A = \text{diag}\left(\begin{bmatrix} 1, \ldots, 10, 0 \end{bmatrix}\right)$ and b is a vector of all ones. Clearly, all but the last components are the same (in general, all components are different), and MINRES-QLP gives the minimum-length solution, whereas MINRES returns a minimum residual solution.

Table IV: MINRES and MINRES-QLP solutions of $Ax \approx e$, where $A = \mathrm{diag}[1, \ldots, 10, 0]$.

MINRES	MINRES-QLP
1.000000000000001	1.000000000000000
0.500000000000001	0.500000000000001
0.333333333333333	0.333333333333333
0.250000000000001	0.250000000000001
0.199999999999999	0.199999999999999
0.166666666666667	0.166666666666667
0.142857142857143	0.142857142857143
0.125000000000000	0.125000000000000
0.111111111111111	0.111111111111111
0.100000000000000	0.100000000000000
2.928968253967685	0.000000000000000

The program produces printed output on file `nout`, if that parameter is positive. This is illustrated below, in which another least-squares problem (Ex. 21 in `minresqlpTestProgram`) is solved: $\min \|x\|$ such that $x \in \arg\min \| \mathrm{diag}\begin{bmatrix} d, 0, 0 \end{bmatrix} x - b\|$, where $d \equiv \begin{bmatrix} \frac{1}{50}, \frac{2}{50}, \ldots, \frac{48}{50} \end{bmatrix}^T$ and $b \equiv \begin{bmatrix} d, 1, 1 \end{bmatrix}^T .* \begin{bmatrix} 50 : -1 : 3, 1, 1 \end{bmatrix}^T$, where $.*$ indicates elementwise multiplication. No preconditioner is applied, and shift $\sigma = 0$.

Notice that the rightmost column of the 39th iteration is marked with "P", which indicates that the program switches from MINRES phase to MINRES-QLP phase since $\mathcal{A}_{39} \approx 1.81 \times 10^7 > trancond = 10^7$. Even though the last line in the output reports that MINRES-QLP has to stop at iteration 46 since $\|x_{47}\| > maxxnorm$, the algorithm appears to be successful because the relative error in x_{46} is merely 2.8×10^{-13}.

```
Enter MINRES-QLP.      Solution of symmetric Ax = b
n = 50                 ||b|| = 6.78E+01          precon = F
itnlim = 200           rtol = 2.22E-16           shift = 0.00E+00
maxxnorm = 1.00E+07    Acondlim = 1.00E+15       trancond = 1.00E+07

iter          x(1)    xnorm    rnorm   Arnorm Compatible     LS  norm(A) cond(A)
  0 0.0000000000E+00 0.00E+00 6.78E+01 3.69E+01   1.00E+00 1.00E+00 0.00E+00 1.00E+00
  1 1.7180943901E+00 1.16E+02 2.40E+01 1.09E+01   1.83E-01 6.94E-01 5.44E-01 1.00E+00
  2 3.8644538109E+00 1.53E+02 1.15E+01 4.58E+00   6.82E-02 6.06E-01 6.57E-01 1.70E+00
  3 6.3954779963E+00 1.72E+02 6.51E+00 2.30E+00   3.60E-02 5.37E-01 6.57E-01 2.27E+00
  4 9.2579303917E+00 1.83E+02 4.16E+00 1.29E+00   2.21E-02 4.74E-01 6.57E-01 2.94E+00
  5 1.2389816033E+01 1.90E+02 2.94E+00 7.91E-01   1.52E-02 4.10E-01 6.57E-01 3.74E+00
  6 1.5722893791E+01 1.95E+02 2.28E+00 5.14E-01   1.16E-02 3.43E-01 6.57E-01 4.78E+00
  7 1.9185796048E+01 2.00E+02 1.91E+00 3.50E-01   9.61E-03 2.79E-01 6.57E-01 6.20E+00
  8 2.2706980590E+01 2.03E+02 1.71E+00 2.48E-01   8.47E-03 2.21E-01 6.57E-01 8.20E+00
  9 2.6217158315E+01 2.07E+02 1.59E+00 1.81E-01   7.81E-03 1.73E-01 6.57E-01 1.10E+01

 10 2.9651001936E+01 2.10E+02 1.52E+00 1.36E-01   7.39E-03 1.36E-01 6.57E-01 1.50E+01
 20 4.9405101158E+01 2.71E+02 1.41E+00 1.08E-02   5.76E-03 1.16E-02 6.57E-01 1.92E+02
 30 4.9999971981E+01 3.22E+02 1.41E+00 6.37E-05   5.06E-03 6.86E-05 6.57E-01 1.18E+04
 39 5.0000000000E+01 3.53E+02 1.41E+00 2.09E-08   4.72E-03 2.25E-08 6.57E-01 1.81E+07 P

 40 5.0000000000E+01 3.56E+02 1.41E+00 6.60E-09   4.69E-03 7.10E-09 6.57E-01 5.29E+07
 45 5.0000000000E+01 3.76E+05 1.41E+00 5.53E-07   5.73E-06 5.95E-07 6.57E-01 2.01E+11
 46 5.0000000000E+01 2.07E+02 1.41E+00 5.86E-07   5.73E-06 6.30E-07 6.57E-01 2.01E+11

Exit MINRES-QLP.      istop = 12               itn = 46
Exit MINRES-QLP.      Anorm = 6.5701E-01       Acond = 2.0123E+11
Exit MINRES-QLP.      rnorm = 1.4142E+00       Arnorm = 1.2149E-05
Exit MINRES-QLP.      xnorm = 2.0717E+02
Exit MINRES-QLP.      xnorm has exceeded maxxnorm or will exceed it next iteration.
```

The items printed at the kth iteration are listed and explained in in the source code. For simplicity we assumed below with no preconditioner; when there is one, we simply replace $\bar{A}$ and r_k, respectively, with $\tilde{A}$ and $\bar{r}_k$ as defined in Section 3 and Algorithm 1.

Table V: Items printed at the kth iteration.

Label	Definition
iter	The iteration number k. Results are always printed for the first 10 iterations and the last. Intermediate results are printed every 10th iteration.

Table V: Items printed at the kth iteration (continued).

Label	Definition
$x(1)$	The value of the first element of the approximate solution x_k.
xnorm	$\|x_k\|$.
rnorm	$\|r_k\|$.
Arnorm	$\|\bar{A}r_k\|$.
Compatible	A dimensionless quantity that should converge to zero if and only if $\bar{A}x = b$ is compatible. It is an estimate of $\|r_k\|/(\|\bar{A}\|\|x_k\| + \|b\|)$, which decreases monotonically.
LS	A dimensionless quantity that should converge to zero if and only if the optimum r_k is nonzero. It is an estimate of $\|\bar{A}r_k\|/(\|\bar{A}\|\|r_k\|)$, which is usually not monotonic.
norm(A)	A monotonically increasing underestimate of $\|\bar{A}\|$.
cond(A)	A monotonically increasing underestimate of $\text{cond}(\bar{A})$.

The integer output `istop` takes an initial value of 0; when the program stops, it takes a positive integer value between 1 to 14 inclusive to signify one of the termination conditions in Table VI. We note that if `istop` > 7, the final x may or may not be an acceptable solution. On the contrary, when `istop` $\leq$ 7, we can be sure x_k is a good or even an excellent approximate solution of a given problem.

Table VI: Termination conditions in MINRES-QLP.

istop	Termination Conditions
1	$\beta_{k+1} < \varepsilon$. Iteration k is the final Lanczos step. Rarely occurs.
2	$\beta_2 = 0$ in the Lanczos iteration; i.e. the second Lanczos vector is zero. This means the right-hand-side is very special. If there is no preconditioner, b is an eigenvector of A.

Table VI: Termination conditions in MINRES-QLP (continued).

istop	**Termination Conditions**
	Otherwise (if *precon* is true), let $My = b$. If shift is zero, y is a solution of the generalized eigenvalue problem $Ay = \lambda My$, with $\lambda = \alpha_1$ from the Lanczos vectors. In general, $(A - \sigma I)x = b$ has the solution $x = (1/\alpha_1)y$, where $My = b$.
3	$b = 0$, so the exact solution is $x = 0$. No iterations were performed.
4	$\|\bar{r}\|$ appears to be less than the value $rtol(\|\bar{A}\|\|\bar{x}\| + \|b\|)$; x should be an acceptable solution of $\bar{A}x = b$.
5	$\|\bar{r}\|$ appears to be less than the value $\varepsilon(\|\bar{A}\|\|\bar{x}\| + \|b\|)$. This means that the solution is as accurate as seems reasonable on this machine.
6	$\|\bar{A}\bar{r}\|$ appears to be less than the value $rtol\|\bar{A}\|\|\bar{r}\|$; x should be an acceptable least-squares solution.
7	$\|\bar{A}\bar{r}\|$ appears to be less than the value $\varepsilon\|\bar{A}\|\|\bar{r}\|$. This means that the least-squares solution is as accurate as seems reasonable on this machine.
8	The iteration limit was reached before convergence.
9	The matrix defined by `Aprod` does not appear to be symmetric. For certain vectors $y = Av$ and $r = Ay$, the products $y^T y$ and $r^T v$ differ significantly.
10	The matrix defined by `Msolve` does not appear to be symmetric. For vectors satisfying $My = v$ and $Mr = y$, the products $y^T y$ and $r^T v$ differ significantly.
11	An inner product of the form $x^T M^{-1} x$ was not positive, so the preconditioner M does not appear to be positive definite.

Table VI: Termination conditions in MINRES-QLP (continued).

istop	Termination Conditions
12	$\|x\|$ has exceeded *maxxnorm* or will exceed it next iteration.
13	$\mathrm{cond}(\bar{A})$ has exceeded *Acondlim* or $0.1/\varepsilon$, so $\bar{A}$ must be very ill-conditioned.
14	$|\gamma_k^{(4)}| < \varepsilon$. This is probably a least-squares problem but residual norms have not satisfied NRBE conditions.

6. AVAILABILITY

Implementations of MINRES-QLP are available in FORTRAN 90 and MATLAB 7.8 from the Systems Optimization Laboratory, Stanford University [SOL], or the first author's homepage http://home.uchicago.edu/sctchoi/ under the terms of the OSI Common Public License (CPL) [OSI-CPL] or the BSD License [BSD].

REFERENCES

BLAS. Basic Linear Algebra Subprograms. http://www.netlib.org/blas/.

BSD. The BSD License. http://www.opensource.org/licenses/bsd-license.php.

BURKARDT, J. 2012. fortran90 software. http://people.sc.fsu.edu/~jburkardt/.

CALVETTI, D., LEWIS, B., AND REICHEL, L. 2000. An L-curve for the MINRES method. In *Proceedings of SPIE*. Vol. 4116. 385–395.

CHIVERS, I. AND SLEIGHTHOLME, J. 2006. *Introduction to Programming with Fortran: With Coverage of Fortran 90, 95, 2003, and 77*. Springer-Verlag, London, UK.

CHOI, S.-C. T. 2006. Iterative Methods for Singular Linear Equations and Least-Squares Problems. Ph.D. thesis, ICME, Stanford University, CA.

CHOI, S.-C. T., PAIGE, C. C., AND SAUNDERS, M. A. 2011. MINRES-QLP: A Krylov subspace method for indefinite or singular symmetric systems. *SIAM J. Sci. Comput. 33,* 4, 1810–1836.

DAVIS, T. A. AND HU, Y. 2011. The University of Florida sparse matrix collection. *ACM Trans. Math. Software 38,* 1, 1:1–1:25.

DONGARRA, J., EIJKHOUT, V., AND KALHAN, A. 1995. Reverse communication interface for linear algebra templates for iterative methods. Report UT-CS-95-291, University of Tennessee, TN.

FONG, D. C.-L. 2011. Minimum-Residual Methods for Sparse Least-Squares Using Golub-Kahan Bidiagonalization. Ph.D. thesis, ICME, Stanford University, CA.

FONG, D. C.-L. AND SAUNDERS, M. A. 2011. LSMR: An iterative algorithm for sparse least-squares problems. *SIAM J. Sci. Comput. 33*, 2950–2971.

FONG, D. C.-L. AND SAUNDERS, M. A. 2012. CG versus MINRES: An empirical comparison. *SQU Journal for Science 17,* 1, 44–62.

FOSTER, L. 2009. San Jose State University Singular Matrix Database. `http://www.math.sjsu.edu/singular/matrices/`.

FREUND, R. W. AND NACHTIGAL, N. M. 1994. A new Krylov-subspace method for symmetric indefinite linear systems. In *Proceedings of the 14th IMACS World Congress on Computational and Applied Mathematics*, W. F. Ames, Ed. IMACS, 1253–1256.

GOLUB, G. H. AND VAN LOAN, C. F. 1996. *Matrix Computations*, 3rd ed. Johns Hopkins University Press, Baltimore, MD.

HANSEN, P. C. 1998. *Rank-deficient and Discrete Ill-posed Problems*. SIAM Monographs on Mathematical Modeling and Computation. SIAM, Philadelphia, PA.

HESTENES, M. R. AND STIEFEL, E. 1952. Methods of conjugate gradients for solving linear systems. *J. Research Nat. Bur. Standards 49*, 409–436.

LANCZOS, C. 1950. An iteration method for the solution of the eigenvalue problem of linear differential and integral operators. *J. Research Nat. Bur. Standards 45*, 255–282.

O'BOYLE, N. M., TENDERHOLT, A. L., AND LANGNER, K. M. 2008. cclib: A library for package-independent computational chemistry algorithms. *J. Comput. Chem. 29,* 5, 839–845.

OLIVEIRA, S. AND STEWART, D. 2006. *Writing Scientific Software: A Guide to Good Style*. Cambridge University Press, Cambridge, UK.

OSI-CPL. Open Source Initiative (OSI) - Common Public License (CPL). `http://www.opensource.org/licenses/cpl1.0.php`.

PAIGE, C. C. 1976. Error analysis of the Lanczos algorithm for tridiagonalizing a symmetric matrix. *J. Inst. Math. Appl. 18,* 3, 341–349.

PAIGE, C. C. AND SAUNDERS, M. A. 1975. Solution of sparse indefinite systems of linear equations. *SIAM J. Numer. Anal. 12,* 4, 617–629.

PAIGE, C. C. AND SAUNDERS, M. A. 1982a. LSQR: an algorithm for sparse linear equations and sparse least squares. *ACM Trans. Math. Software 8,* 1, 43–71.

PAIGE, C. C. AND SAUNDERS, M. A. 1982b. Algorithm 583; LSQR: Sparse linear equations and least-squares problems. *ACM Trans. Math. Software 8,* 2, 195–209.

PAIGE, C. C. AND STRAKOŠ, Z. 2002. Residual and backward error bounds in minimum residual Krylov subspace methods. *SIAM J. Sci. Comput. 23,* 6, 1899–1924.

SAAD, Y. AND SCHULTZ, M. H. 1986. GMRES: a generalized minimal residual algorithm for solving nonsymmetric linear systems. *SIAM J. Sci. Statist. Comput. 7,* 3, 856–869.

SOL. Systems Optimization Laboratory, Stanford University. `www.stanford.edu/group/SOL`.

STEWART, G. W. 1999. The QLP approximation to the singular value decomposition. *SIAM J. Sci. Comput. 20,* 4, 1336–1348.

TREFETHEN, L. N. AND BAU, III, D. 1997. *Numerical Linear Algebra*. SIAM, Philadelphia, PA.

reproduce, prepare derivative works, distribute copies to the public, and perform publicly and display publicly, by or on behalf of the Government.